



UNSW
AUSTRALIA

Overtone filtration techniques for music spectrum analysis (COMP4121)

There exists many tools for spectral analysis of audio using algorithms such as the Fast Fourier transform and the discrete cosine transform, yet they are not very accurate when used to transcribe notes from an audio input.

This paper explores two approaches to supplementing these spectral analytic methods, by improve their results. In particular, we explore methods involving markov chains and probabilistic filtering of overtones.

Contents

1	Previous work and acknowledgments	3
1.1	Attribution of source code	3
1.2	DFT	3
1.3	Converting frequencies into notes	3
1.4	Window functions	4
2	Problems with naive DCT based transcriptions	4
2.1	Harmonic and Inharmonic interference	4
2.2	Inaccuracy when multiple notes are played	5
3	A digression into why modelling as a stochastic process would be ineffective	5
4	An iterative method for filtering out interfering overtones	6
4.1	Aim of method	6
4.2	Overtone Distributions	7
4.3	How do we find O ?	7
4.4	Noise	8
4.5	Initial distribution	9
4.6	Approximators	9
4.6.1	Echo approximator	9
4.6.2	Average of history	9
5	Results and further improvement	9
5.1	Results	9
5.2	Ideas for further improvement	9
	Appendices	10
A	Transcriptions	10
A.1	The Well-Tempered Clavier (pedal)	10
A.1.1	With our augmentations	10
A.1.2	Without our augmentations	11
A.2	The Well-Tempered Clavier (no-pedal)	12
A.2.1	With our augmentations	12
A.2.2	Without our augmentations	14
A.3	Polyphonic	15
A.3.1	With our augmentations	15
A.3.2	Without our augmentations	16
6	Bibliography	17

1 Previous work and acknowledgments

1.1 Attribution of source code

The work in this paper builds upon already existing work on the effectiveness of spectral analysis techniques for music transcription, specifically that of a previous taker of this course (Gregory Omelaenko). When constructing experiments to gauge the effectiveness of the techniques described in this paper, source code originally written by the author of the reference paper, but heavily modified (ported to a different language, different execution environment and to use different libraries) for the purpose of this paper was reused.

The code that was reused includes the implementations of the DCT/FFT algorithms, conversion from frequencies to notes, the window functions, the overlapping of samples and the code that combines these into an output result. The rest of this section skims over the sections of said paper of which are particularly relevant to this one, for the sake of locality.

1.2 DFT

The Discrete Fourier Transform (henceforth referred to as the DFT, or FFT) is an algorithm that decomposes a discrete signal into a sum of sine waves, encoding their frequency, amplitude and phase. It can be calculated in $O(n \log n)$. The DFT $X_0 \dots X_{N-1}$ of sequence $x_0 \dots x_{N-1}$ is defined as follows:

$$X_k = \sum_{n=0}^{N-1} x_n e^{\frac{-2\pi i k n}{N}}$$

The Discrete Cosine Transform is a similar transform, but one that only uses real numbers. It can also be calculated in $O(n \log n)$. Whilst there are many variations of it, all further references to the DCT refer to the DCT-II, which is defined as follows:

$$C_k = \sum_{n=0}^{N-1} x_n \cos \left(\frac{\pi}{N} \left(n + \frac{1}{2} \right) k \right)$$

1.3 Converting frequencies into notes

When a note is played, the resultant sound is the summation of a sequence of sine waves, with coefficients and frequencies varying by the note's timbre. The lowest frequency that is always present when the note is played is called the *fundamental frequency*, and is considered the frequency of the note. Any higher frequency associated with the note is called an *overtone*, the overtones and the fundamental are together called the *partials* of the note. From this, we can associate every note in the western musical scale with a frequency by associating it with its *fundamental frequency*.

The association works as follows, we number each note j by the number of semitones above A0 it is, $\#n(j)$. Then the frequency f_j associated with note j is:

$$f_j = 440 \cdot 2^{\frac{\#n(j)}{12} - 4}$$

We can use this to extract notes from our DCT/DFT. This can be further enhanced by associating a 'bucket' of frequencies with each note, by partitioning the frequency spectrum by their (logarithmically) closest associated frequency, allowing the entire FFT/DCT output to be used when classifying notes.

1.4 Window functions

Since the signal is usually not periodic in the number of samples in the DFT window, spectral leakage occurs. A simple way to combat this is via the application of a window function, which gives higher weighting to samples in the middle of the window. A window function is any function satisfying:

$$\begin{aligned} w : \mathbb{Z} \times \mathbb{N}_{>0} &\rightarrow [-1, 1] \\ i \notin [0, n) &\Rightarrow w(i, n) = 0 \end{aligned}$$

Application of a window function is simple: to apply the window function to a sample x_i in a window of size n , you simply multiply x_i by $w(i, n)$. The window function used in the reference analyser, and also used here is the Blackman-Harris window:

$$w_{\text{BH}}(i, n) = \begin{cases} 0.422323 - 0.49755 \cos\left(\frac{2\pi i}{n}\right) + 0.7922 \cos\left(\frac{4\pi i}{n}\right) & \text{if } i \in [0, n) \\ 0 & \text{otherwise} \end{cases}$$

Overlapping samples are also employed to lower the chance of missing a note due to it being shorter than the DFT window, which is a problem made only worse by the use of a window function since some sections of the audio would be given very little weighting.

2 Problems with naive DCT based transcriptions

2.1 Harmonic and Inharmonic interference

One of the major issues observed in the operation of the spectral analysis transcriber is that it has considerable difficulty distinguishing notes from their octave harmonics. Whenever a note is played, it's fundamental frequency is sounded, but typically the harmonics are also sounded as well. As a result, not all tones that appear in the FFT/DCT correspond to an actual note being played, so these overtones must be removed before note analysis is performed. However, some of these harmonic tones may also be the fundamental frequencies of other

notes that are being played! This leaves us with the quite difficult task of figuring out how much of the volume of each tone is caused by a note below it.

The author of the reference transcriber mistakenly assumes that a note can only interfere with, or be interfered by, a note whose fundamental frequency differs from it by a power of two. As a result of this, their strategy to reduce interference only worries about notes that differ by an integer number of octaves from each other, which could cause problems where the interference is on some non power of two harmonic. As an example, the third harmonic of a note might not be recognized as a harmonic of that note. If this tone has sufficient volume, a *just-perfect* fifth chord could be detected, where none occurred.

However, not even accounting for all harmonic frequencies would be sufficient to solve this problem. Not only can a note create harmonic overtones, it can also, in some cases, cause inharmonic overtones. Certain ways of playing instruments, and some instruments in general cause inharmonic tones, eg a string instrument being plucked, or most percussion instruments being played. This means that any solution that operates by filtering out overtones would have to account for possible interference between notes of any two frequencies, to be generally applicable.

2.2 Inaccuracy when multiple notes are played

Whilst the methods used by the reference transcriber are somewhat effective at identifying up to two notes, having more than that will cause entire notes to be missed. This is probably due to the excessive filtering needed to extract notes from simpler samples, which would result in correct notes being discarded as noise in more complex pieces.

3 A digression into why modelling as a stochastic process would be ineffective

A seemingly-simple method that comes quickly to mind when encountering this problem is to use information about how notes map to tones and what sequences of notes are expected to be common in the audio sample to create a Hidden Markov Model, and apply the viterbi algorithm to augment the results.

The biggest reason why this would not work is state-space explosion. The case where we have a very simple melody does not really require any augmentation, since the methods used by the spectral analyser are by and large effective enough for these cases. What we really care about are cases with multiple simultaneous notes being played. If we model the markov chain as states are the individual notes being played, we cannot really use it on samples where multiple notes are being played.

A way to get around this would be to consider combinations of notes as our states. This gives us a state space of size $|S| = 2^N$, where N is the number of notes we account for (which in practice would be at least 88, which is the number of keys on a piano). This is already infeasible, since Viterbi requires

$O(|S|^2 + |S|O)$ space, and $O(|S|^2 \times T)$ time to run. Even if we were to limit this by only allowing up to k notes to be pressed at a time, we would still end up with $|S| = C_k^n$, which still grows very fast in k .

This problem becomes even worse when polyphonic music is considered. Monophonic music is music that only has one 'voice', in which case there are very strong structural ties between consecutive notes and chords in, which case Viterbi could be somewhat useful. Polyphonic music has multiple voices, and the Viterbi algorithm is incapable of leveraging knowledge about individual voices to reason about their composition.

Modelling a piece of music as a stochastic process would require some way of partitioning notes into their corresponding voices to succeed, which is incredibly non-trivial, as the notes of two different voices can stay very close to each other. If this problem were to be solved, the problem of state space explosion would be just-barely avoidable, since each voice could realistically produce 5 or less notes simultaneously.

4 An iterative method for filtering out interfering overtones

4.1 Aim of method

We propose a radical simplification to our model to make it more tractable, one that is somewhat justifiable, but not necessarily always true. We shall assume that all notes of an instrument are approximately the same, just frequency shifted by the ratios of their fundamental frequencies. We shall also assume that there is only one instrument in the playback (though this is a much simpler assumption).

This implies that all notes have the same overtone distribution, which greatly simplifies the process of filtering out the overtones of a note. In practice, this is not true, the distribution is different for different notes of the same instrument, and even different instances of the same note. There are a lot of reasons behind this, most of them related to psycho-acoustics, some of them related to constructive and destructive interference, how the human ear responds to different frequencies, and the compensation methods employed by the instruments to counteract these effects. The assumption is still somewhat justifiable though, since the distribution of these overtones partially determines the timbre of the instrument, and since instruments typically aim to sound uniform in timbre across their musical ranges.

We propose an iterative method for approximating the distribution of overtones for an instrument given a series of observed volumes of tones, and to use it to better isolate notes from the output of the DCT/FFT. This method takes advantage of the fact that notes can only interfere with notes above them, not notes below, and attempts to figure out the overtone distribution from bottom-up.

4.2 Overtone Distributions

We assume, like the spectral analyser does, that there are only 137 different pitches to distinguish, and operate on a vector of 137 non-negative real numbers representing the volumes of each tone as detected in the sample. We define this vector as:

$$V \in (\mathbb{R}_0^+)^{137}$$

$$V_i = \text{Volume of note } i \text{ semitones above } A0$$

We then define a vector of overtone ratios O , henceforth also referred to as an overtone distribution, and a vector of determined notes N as follows:

$$O \in (\mathbb{R}_0^+)^{136}$$

$$N \in (\mathbb{R}_0^+)^{137}$$

$$O_i = \text{ratio of volume of Fundamental Frequency}$$

$$\text{to volume of overtone } i+1 \text{ semitones above}$$

$$V_i = N_i + \sum_{k=0}^{i-1} N_k O_{i-1-k}$$

Whilst negative overtone ratios are possible, due to destructive interference, we assume that they are anomalous, and model O as only containing non-negative values. From this, it follows that:

$$N_i = V_i - \sum_{k=0}^{i-1} N_k O_{i-1-k}$$

It becomes apparent that $V_0 = N_0$, and that we can calculate the value of N by iteratively solving for each successive element in N , and subtracting the overtone distribution for that term. Which gives us the following iterative algorithm for finding N given O and V (NOTE: in the pseudocode below, sequences such as 1..2 include the start but not the end):

```
for i in 0 .. 137:
    N[i] := V[i]
    for j in 0 .. 136-i:
        V[i+j+1] := V[i+j+1] - N[i] * O[j]
```

4.3 How do we find O ?

We need to actually calculate O before we can use it to remove overtones.

One way of doing so is having a single known note played a couple times at the beginning, so as to get a rough estimate of the instrument's overtone distribution.

We can use feedback from N to iteratively calculate better approximations of O . We have to filter out the noise from N to figure out which notes were actually

present, and which notes were noise. We set the notes that were determined to be noise to 0 in N . We then calculate O' to minimize the difference between the V and V' , where V' is calculated from O' and the filtered N .

Calculating O' can be expensive, however we note that we care a lot more about smaller overtone frequency ratios than larger ones, since their volume ratios tend to be higher. Hence, we can make our calculation of O' greedy without too much loss in accuracy, by trying to maximise smaller frequency ratios first. This allows the algorithm to be more efficient.

We then feed O' into whatever algorithm we are using to approximate O , which will give us the next approximation.

Our new algorithm for determining N , given V and an approximation of O is:

```
O := approximator.get()
V_t := V.clone()
for i in 0 .. 137:
    if V_t[i] < noise_threshold:
        N[i] := 0
    else:
        N[i] := V_t[i]
        for j in 0 .. 136-i:
            V[i+j+1] := max(0, V[i+j+1] - N[i] * O[j])
O_t := new double[136]
V_t := V.clone()
for i in 0 .. 137:
    V_t[i] := V_t[i] - N[i]
for i in 0 .. 136:
    O_t[i] = 0
    j := 0
    while j+i+1 < 137 && (N[j] == 0 || N[i+j+1] == 0):
        j := j + 1
    if j+i+1 != 137:
        O_t[i] = V_t[j+i+1]/N[j]
        for j in 0 .. 137-i-1:
            V_t[i+j+1] := max(0, V_t[i+j+1] - N[j] * O_t[i])
approximator.feed(O_t)
```

4.4 Noise

For the purpose of detecting noise, we shall reuse the noise filtration system used in the Spectral analyser. We simply consider the volume of the loudest notes of this sample v and the threshold of the previous sample t' and set the noise threshold $t = \max(sv, t'd)$ for some constants $0 \leq s, d \leq 1$. For our experiments, we set $d = 0.7, s = 0.4$

4.5 Initial distribution

We now have to choose an initial distribution. Whilst simply choosing zero would probably be sufficient, we make the observation that typically, the overtones of a note are its harmonics. With this in mind, we choose an initial vector O that is zero for all elements O_i except where $2^{i/12}$ is the closest value to some integer $1 < k < 7$, in which case $O_i = 1/(k * k)$.

4.6 Approximators

Now we get on to the subject of how we design our approximators. An approximator is some opaque object that can be queried for an overtone vector determined by the overtone vectors that were fed into it so far.

4.6.1 Echo approximator

The simplest approximator possible would be the Echo approximator. It basically returns the last vector that was fed in to it. For obvious reasons, this would be an awful approximation.

4.6.2 Average of history

A better approximator would be one that returns a mean of some number of previous overtone vectors. In particular, it would average the last n vectors, which are initially just the initial vector n times. This is the approximator we use.

5 Results and further improvement

5.1 Results

In practice, the techniques shown in this paper do not greatly increase the quality of the transcription. There are a few areas where this algorithm creates a slightly more accurate, and it is never noticeably worse than the baseline, so it can tentatively be called an improvement over the naive spectral analyser. Comparisons between transcriptions with and without this algorithm can be found in the appendix.

It is, however, noteworthy that when the outputs of the analysis were fed back into a synthesizer (synth.cc in the source code), the songs that were transcribed were mostly recognizable, and in some sections fairly correct. This implies that whilst it is not sufficient for automated audio transcription, it might be still be useful as a guide for manual transcription.

5.2 Ideas for further improvement

Using linear programming to calculate the optimal O' , instead of greedily calculating a lower-term biased O' , would probably improve the results in heavily

pedaled pieces. I believe in retrospect that the greedy calculation of O' , whilst favouring the frequencies that are closer to the fundamental, also greatly favours the lower frequency of notes in general. This causes the base of any octave chord to dominate the higher part, which causes inaccuracies for certain sequences of notes (eg. arpeggios being played with pedal).

Another interesting idea for further improvement is using maximum likelihood estimation (or some analogue of it) for choosing an approximation for O based on past O 's, instead of simply averaging.

One final, if not very well fleshed out potential avenue for advancement lies again in the Viterbi algorithm. Were it possible to collapse the state space heavily, or find some way of effectively partitioning the observed tones into their respective voices, Viterbi might be feasible.

Appendices

A Transcriptions

Note: The numbers represent the time elapsed in $\frac{1}{32}$ -seconds.

A.1 The Well-Tempered Clavier (pedal)

This is an excerpt from the attempted transcription of the beginning of a pedaled recording of Bach's 'The Well-Tempered Clavier'.

A.1.1 With our augmentations

109: C4 down B3 down C#4 down	510: G4 down C4 up
110: B3 up C#4 up	517: C5 down
128: C4 up E4 up	528: C5 up
136: C5 down	556: C4 down C5 down
149: C5 up	560: G4 up C5 up
156: C5 down	571: A4 down G4 down
169: C4 down	576: A4 up
170: C5 up	578: D#5 up C4 up E4 up
172: C5 down	585: A5 down E5 up
177: E4 down C5 up	591: A4 down
179: E5 down C5 down	597: A4 up
180: C5 up E5 up	612: C4 down C5 down
182: C4 up E4 up	617: C5 up
190: C5 down	624: A5 up
201: C5 up	625: A4 down
210: C5 down	629: A4 up
216: C5 up	632: C4 up E4 up
223: C4 down	639: A5 down

245: D4 up C4 up	646: A4 down C4 down
259: A4 down C4 down	650: C4 up A4 up
261: F5 down C4 up	653: A5 up
264: D4 up A4 up F5 up	660: A5 down
279: C4 down C5 down	667: C4 down
284: C5 up	669: E5 up A5 up
287: C4 up	675: C4 up
292: C4 down	678: C4 down
299: D4 up C4 up	680: C4 up
310: C4 down	722: C4 down
314: A4 down	727: D4 down C#4 down
319: C#4 down	728: C#4 up
320: D5 down C4 up C#4 up D4 up A4 up F5 up	734: C4 up
334: C4 down	740: C4 down
335: B3 down C4 up D5 up F5 up	741: C4 up
353: B3 up	745: C4 down
355: D5 down B3 down D4 up	747: C4 up D4 up
356: B3 up	775: B3 down
365: B3 down	789: G4 down B3 up D4 up
368: B3 up D5 up F5 up	794: B3 down
388: B3 down	795: D5 down B3 up
410: B3 up D4 up	802: B3 down
412: B3 down	807: B3 up
420: B3 up D4 up F5 up	828: B3 down G4 up
442: C4 down B3 down	848: B3 up
443: B3 up	855: B3 down
455: C4 up E4 up E5 up	860: B3 up D5 up
463: C5 down	882: B3 down
474: C5 up	888: C4 down
483: C5 down	897: C4 up
492: C5 up	902: B3 up
496: C4 down	907: B3 down

A.1.2 Without our augmentations

109: C4 down B3 down C#4 down	300: D5 down
110: C#4 up B3 up	307: F5 down
111: C5 down	310: C4 down D4 down
117: C5 up	315: A4 down
121: E4 down	318: A4 up
122: C4 up	320: C4 up D4 up F5 up
124: C4 down E5 down	328: F5 down
127: E5 up	335: B3 down D5 up F5 up
128: E4 up C4 up	342: D4 down
129: G4 down	349: G4 down
136: C5 down	351: D4 up

142: E5 down	353: B3 up
149: E5 up C5 up	355: D5 down
156: C5 down	363: F5 down
169: C4 down	366: B3 down
172: G4 up	367: D5 up
177: E4 down	368: B3 up F5 up
179: E5 down	376: D5 down
181: E5 up	382: F5 down
182: C4 up E4 up C5 up	384: F5 up
183: G4 down	386: D5 up
190: C5 down	388: B3 down
196: E5 down	390: G4 up
201: E5 up C5 up	394: D4 down
211: C5 down	397: D5 down
213: C5 up	398: D5 up
219: E5 down	402: G4 down
224: C4 down	404: B3 up
228: G4 up	405: D4 up
229: C5 down	406: B3 down
230: C5 up E5 up	407: D4 down
231: D4 down	408: D4 up
234: C4 up	409: D5 down
241: C4 down	410: B3 up
243: C4 up	415: B3 down
245: D4 up	419: D5 up
246: D5 down	420: B3 up
253: F5 down	429: D5 down
254: F5 up	435: F5 down
257: D4 down	439: F5 up
259: A4 down	440: D5 up
261: F5 down	442: C4 down B3 down
263: F5 up	443: B3 up
264: A4 up D4 up	445: C5 down
279: C4 down C5 down D5 up	447: C5 up G4 up
285: C5 up	448: E4 down
286: D4 down	452: E5 down
287: C4 up	455: C4 up E4 up E5 up
292: C4 down	456: G4 down
299: D4 up C4 up	463: C5 down

A.2 The Well-Tempered Clavier (no-pedal)

This is an excerpt from the attempted transcription of the beginning of a non-pedaled recording of Bach's 'The Well-Tempered Clavier'.

A.2.1 With our augmentations

27: C4 down	186: F5 up
29: G5 down	190: A4 down A#4 down
30: G5 up	191: A#4 up
34: E4 down D#4 down	195: A4 up
35: C4 up D#4 up	199: D5 down
43: G4 down	207: F5 down D5 up
51: C5 down	212: F5 up
56: G4 up	214: C4 down C#4 down G5 down C5 down
57: E4 up C5 up	215: C#4 up C5 up
58: E5 down	218: G5 up
60: E5 up	222: C4 up
63: E4 down	223: D4 down
66: G4 down E4 up	231: A4 down
71: E4 down G4 up	238: D5 down
72: E4 up	240: A4 up
74: C5 down	244: C4 down
81: C5 up	245: D5 up C4 up
82: E5 down	246: F5 down D4 up
90: C4 down G5 down E5 up	249: F5 up
93: G5 up	251: D4 down
97: E4 down D#4 down C4 up	252: D4 up
98: D#4 up	253: A4 down
106: G4 down	257: A4 up
114: C5 down G4 up	261: D5 down
116: C5 up	268: F5 down
120: E4 up	269: D5 up
121: E5 down	273: F5 up
123: E5 up	276: B3 down A#3 down
124: E4 down	278: A#3 up
128: E4 up	285: D4 down B3 up
129: G4 down	288: B3 down
132: G4 up	292: G4 down
134: E4 down	300: D5 down
135: E4 up	302: D4 up
136: C5 down	305: D4 down G4 up
143: C5 up	307: F5 down
144: E5 down	309: D5 up
150: E5 up	310: D4 up
151: C4 down C#4 down C5 down G5 down	311: F5 up
152: C#4 up C5 up G5 up	313: D4 down
154: G5 down	314: D4 up
155: G5 up	315: G4 down B3 up
160: D4 down C#4 down	320: G4 up
161: C4 up C#4 up	322: D5 down
168: A4 down	331: F5 down
170: D4 up	332: D5 up

176: D5 down A4 up
 178: D4 down
 181: A4 down
 183: F5 down D4 up A4 up D5 up

338: F5 up
 339: B3 down
 347: B3 up
 348: D4 down

A.2.2 Without our augmentations

27: C4 down
 34: E4 down C4 up
 43: G4 down
 51: C5 down
 53: C5 up
 54: G4 up
 55: C5 down
 56: C5 up
 57: E4 up
 58: E5 down
 60: E5 up
 64: E4 down
 65: E4 up
 66: G4 down
 69: G4 up
 74: C5 down
 81: C5 up
 82: E5 down
 89: E5 up
 90: C4 down G5 down
 92: G5 up
 96: C4 up
 97: E4 down
 106: G4 down
 114: G4 up
 120: E4 up
 121: E5 down
 123: E5 up
 125: E4 down
 128: E4 up
 129: G4 down
 132: G4 up
 136: C5 down
 138: C5 up
 140: C5 down
 143: C5 up
 144: E5 down
 150: E5 up
 151: C4 down C#4 down C5 down G5 down

183: F5 down D5 up
 186: F5 up
 190: A4 down A#4 down
 191: A#4 up
 194: A4 up
 199: D5 down
 206: D5 up
 207: F5 down
 212: F5 up
 214: C4 down C#4 down C5 down G5 down
 215: C#4 up
 216: C5 up
 218: G5 up
 222: C4 up
 223: D4 down
 231: A4 down
 233: D4 up
 238: D5 down
 240: D4 down A4 up
 245: D5 up D4 up
 246: F5 down
 249: F5 up
 253: A4 down
 257: A4 up
 261: D5 down
 267: D5 up
 268: F5 down
 273: F5 up
 276: B3 down
 285: D4 down B3 up
 290: B3 down
 292: G4 down
 297: B3 up
 299: B3 down D4 up
 300: D5 down
 302: B3 up
 304: B3 down
 305: G4 up
 307: F5 down

152: C#4 up G5 up
 153: C5 up
 154: G5 down
 155: G5 up
 160: D4 down C4 up
 168: A4 down
 170: D4 up
 176: D5 down A4 up
 178: D4 down
 181: A4 down
 182: D4 up A4 up

309: D5 up
 311: F5 up
 315: G4 down B3 up
 320: G4 up
 322: D5 down
 331: F5 down
 332: D5 up
 338: F5 up
 339: B3 down
 347: B3 up
 348: D4 down

A.3 Polyphonic

This is an excerpt from the attempted transcription of a polyphonic piece. The piece is 'Bach-Brahms: Chaconne in D minor, performed by Krystian Zimmerman'.

A.3.1 With our augmentations

48: D0 down
 49: G0 down
 51: G0 up
 52: C0 down
 53: C0 up D0 up
 55: C0 down
 56: D0 down C0 up
 58: C0 down
 59: C0 up D0 up
 60: F0 down
 61: D0 down
 62: D0 up F0 up
 64: C0 down
 65: C0 up
 67: D0 down
 68: D0 up
 70: F0 down
 72: F0 up
 73: D0 down C0 down
 74: C0 up D0 up
 75: A3 down A4 down
 76: C5 down
 78: F4 down
 79: C5 up
 84: G#3 down
 86: G#3 up

134: C0 down D0 up
 135: F4 down C0 up
 136: D0 down
 137: D0 up F4 up
 138: D4 up A4 up
 139: C0 down
 140: D0 down
 141: C0 up D0 up
 142: A3 down A4 down F4 down
 144: C5 down
 147: C5 up
 156: A3 up
 160: A3 down
 162: F4 up A3 up
 163: E4 down C0 down D#4 down
 164: C0 up D#4 up A4 up
 165: A#3 down A3 down
 167: A3 up
 179: D4 down
 180: D4 up
 182: D4 down
 183: A#3 up D4 up
 185: A#3 down D4 down
 186: D4 up
 188: D4 down
 190: A4 down

92: E5 down	191: A4 up
96: D4 down E5 up	193: D0 down C0 down G3 down
98: D4 up F4 up	194: C0 up
101: G#3 down F4 down	195: G3 up D0 up
102: G#3 up	197: G3 down
103: D4 down	204: A#3 up
105: F4 up	205: C0 down D0 down
107: D3 down	206: C0 up D0 up E4 up
108: G#3 down D3 up	207: G3 up D4 up
109: C0 down	208: C#2 down D0 down F0 down
110: E5 down D3 down C0 up G#3 up	209: F0 up D0 up C#2 up
111: D3 up E5 up	210: C0 down
112: F4 down	211: D0 down
114: F4 up	212: C0 up D0 up
115: C0 down	216: E4 down
117: C0 up A4 up	223: A4 down
118: F0 down	224: A4 up
119: D0 down A3 up F0 up	226: A4 down
120: D4 up	227: A4 up
121: C0 down	229: A4 down
122: D0 up	231: A4 up
126: D0 down C0 up	238: A4 down
129: D4 down	257: G3 down
132: A4 down	260: A4 up

A.3.2 Without our augmentations

48: D0 down	178: A#3 down
50: D0 up	181: A#3 up
52: D0 down C0 down	192: D4 down
53: D0 up C0 up	193: A#3 down
57: D0 down	194: D0 down A#3 up
58: C0 down	195: D0 up
59: C0 up D0 up	197: G3 down A#3 down
61: F0 down	203: A#3 up
62: F0 up	204: E4 up
64: C0 down	205: C0 down D0 down
65: C0 up	206: D4 up C0 up D0 up
67: D0 down	207: G3 up
68: D0 up	208: C#2 down D0 down F0 down
70: F0 down	209: F0 up D0 up C#2 up
71: F0 up	210: C0 down
73: D0 down	211: D0 down
74: D0 up	212: C0 up D0 up
75: A4 down	216: E4 down
76: A3 down	217: E5 down

82: F4 down	219: E5 up
95: E5 down	250: A4 down
96: E5 up	251: A4 up
98: F4 up	259: G3 down
106: D4 down	260: G3 up
109: G#3 down C0 down	263: G3 down
110: C0 up G#3 up	265: D0 down C0 down
115: C0 down	266: E4 up G3 up D0 up
116: A3 up A4 up	267: C0 up
117: C0 up	271: C0 down
119: D4 up	273: C0 up
121: C0 down	280: E4 down
126: C0 up	306: F4 down E4 up
130: D4 down	335: F4 up
133: A4 down	339: F4 down A#3 down
136: D0 down	348: F4 up
137: D0 up	349: D4 down D5 down A#3 up
138: D4 up A4 up	352: A#3 down
139: C0 down	354: A#3 up
141: C0 up	360: A#3 down
142: A4 down A3 down	361: A#3 up D5 up
148: F4 down	368: D5 down
152: F4 up	371: A#3 down A#4 down
154: F4 down	372: A#3 up D5 up
155: A3 up	374: A#3 down
161: A3 down	375: F4 down
162: F4 up A3 up	376: C0 down D0 down
163: E4 down	377: C0 up D0 up
164: A4 up	379: C0 down
166: A#3 down	381: D4 up A#3 up A#4 up C0 up
168: A#3 up	383: A#4 down

6 Bibliography

1. <https://en.wikipedia.org/wiki/Overtone>
2. https://en.wikipedia.org/wiki/Maximum_likelihood_estimation
3. <https://en.wikipedia.org/wiki/Loudness>
4. [https://en.wikipedia.org/wiki/Pitch_\(music\)](https://en.wikipedia.org/wiki/Pitch_(music))
5. <https://en.wikipedia.org/wiki/Harmonic>
6. [https://en.wikipedia.org/wiki/Transcription_\(music\)#Pitch_detection](https://en.wikipedia.org/wiki/Transcription_(music)#Pitch_detection)
7. Greg's source code for his 4121 project.
8. Greg's report for his 4121 project.

9. <https://en.wikipedia.org/wiki/Inharmonicity>